

# Lecture 34

## Learning

# Complex Distributions

*Sankha Narayan Guria*


*With slides from*

*Armando Solar-Lezama*

# n-gram models

---

The big brown bear scares the children with its roar


$$P(\text{scares} | \text{bear}, \text{brown})$$

Probability of a word depends on the previous  $n$  words

Represented with a table:  $P(w_i | w_{i-1}, w_{i-2}, \dots, w_{i-n})$

Bigger  $n$  makes more accurate, but also more difficult to learn, requires much bigger table

Downsides

- some words require more context than others
- some words carry very little information


# Recurrent models

---

Key idea: let the system figure out how to construct its own context

Now need to learn two interrelated functions

- $P(word_i | context_{i-1})$
- $context_i = f(word_i, context_{i-1})$

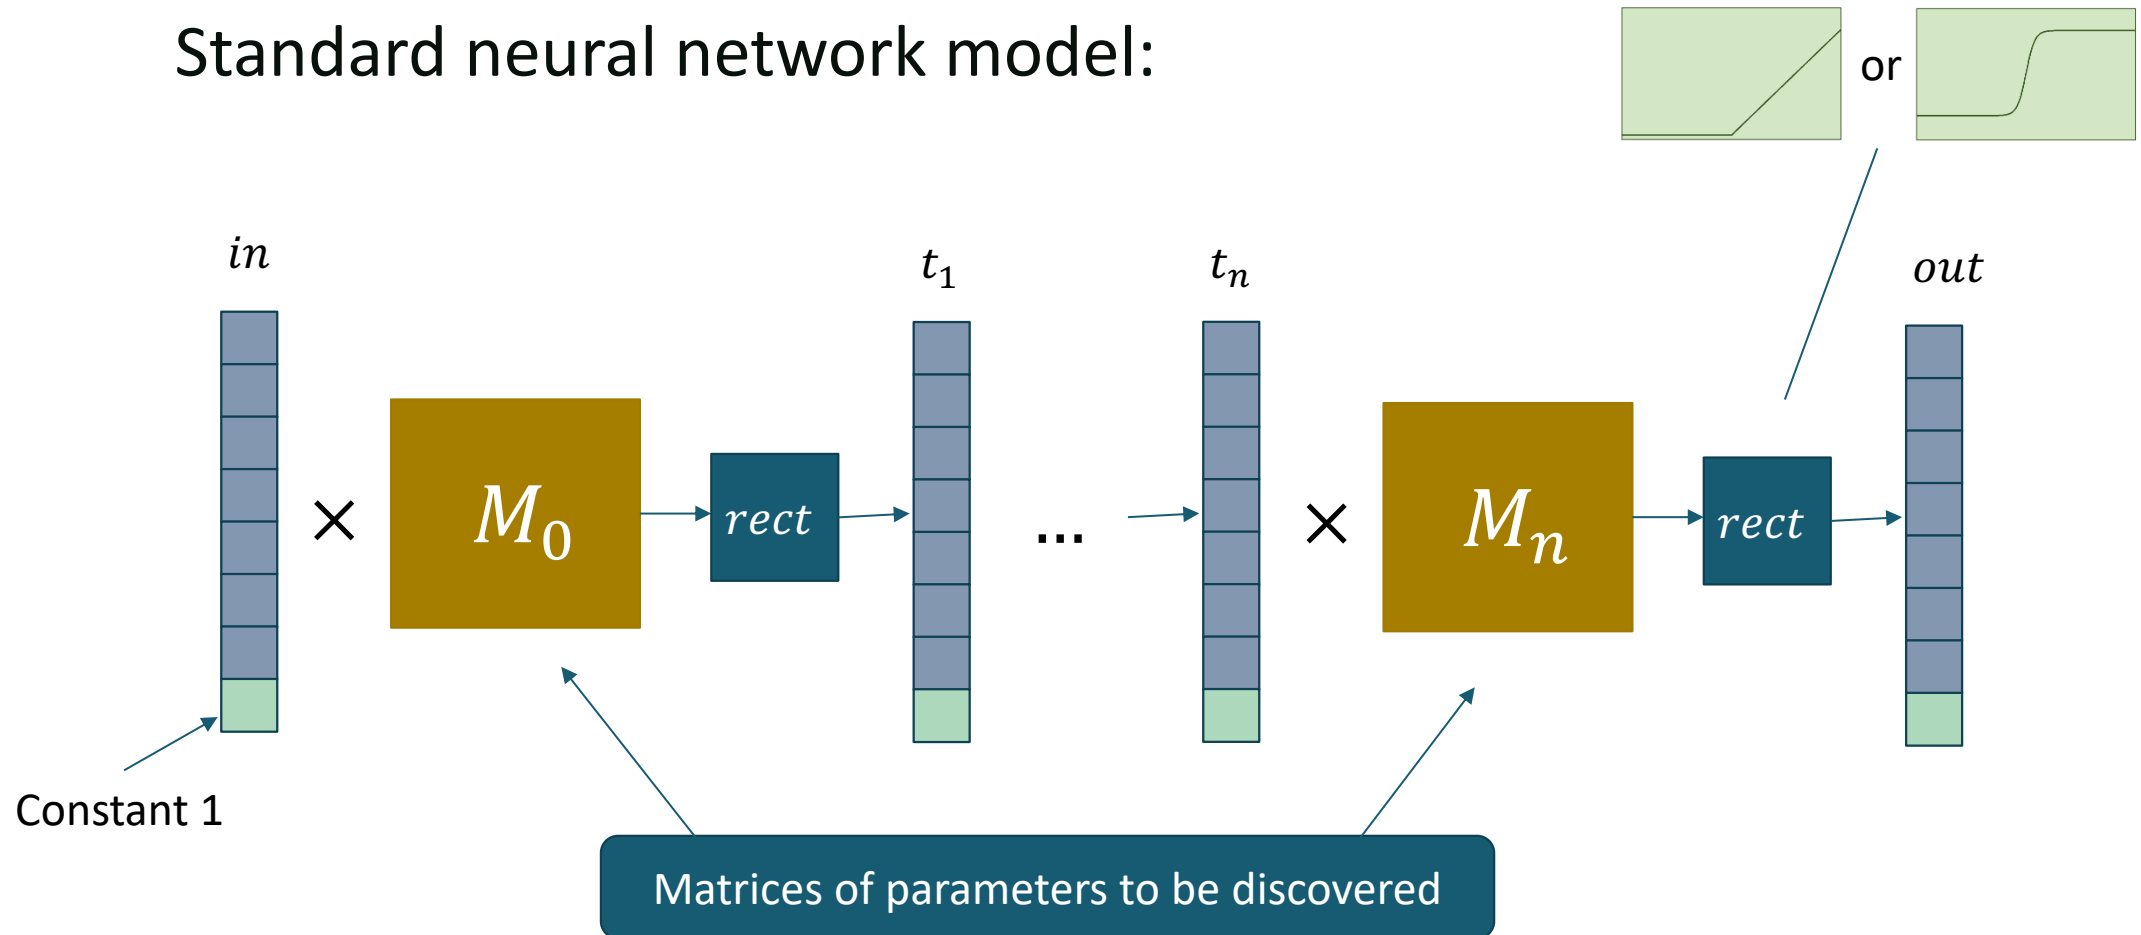


The big brown bear scares the children with its roar

# Neural Networks



Most popular way of representing recurrent models

Standard neural network model:

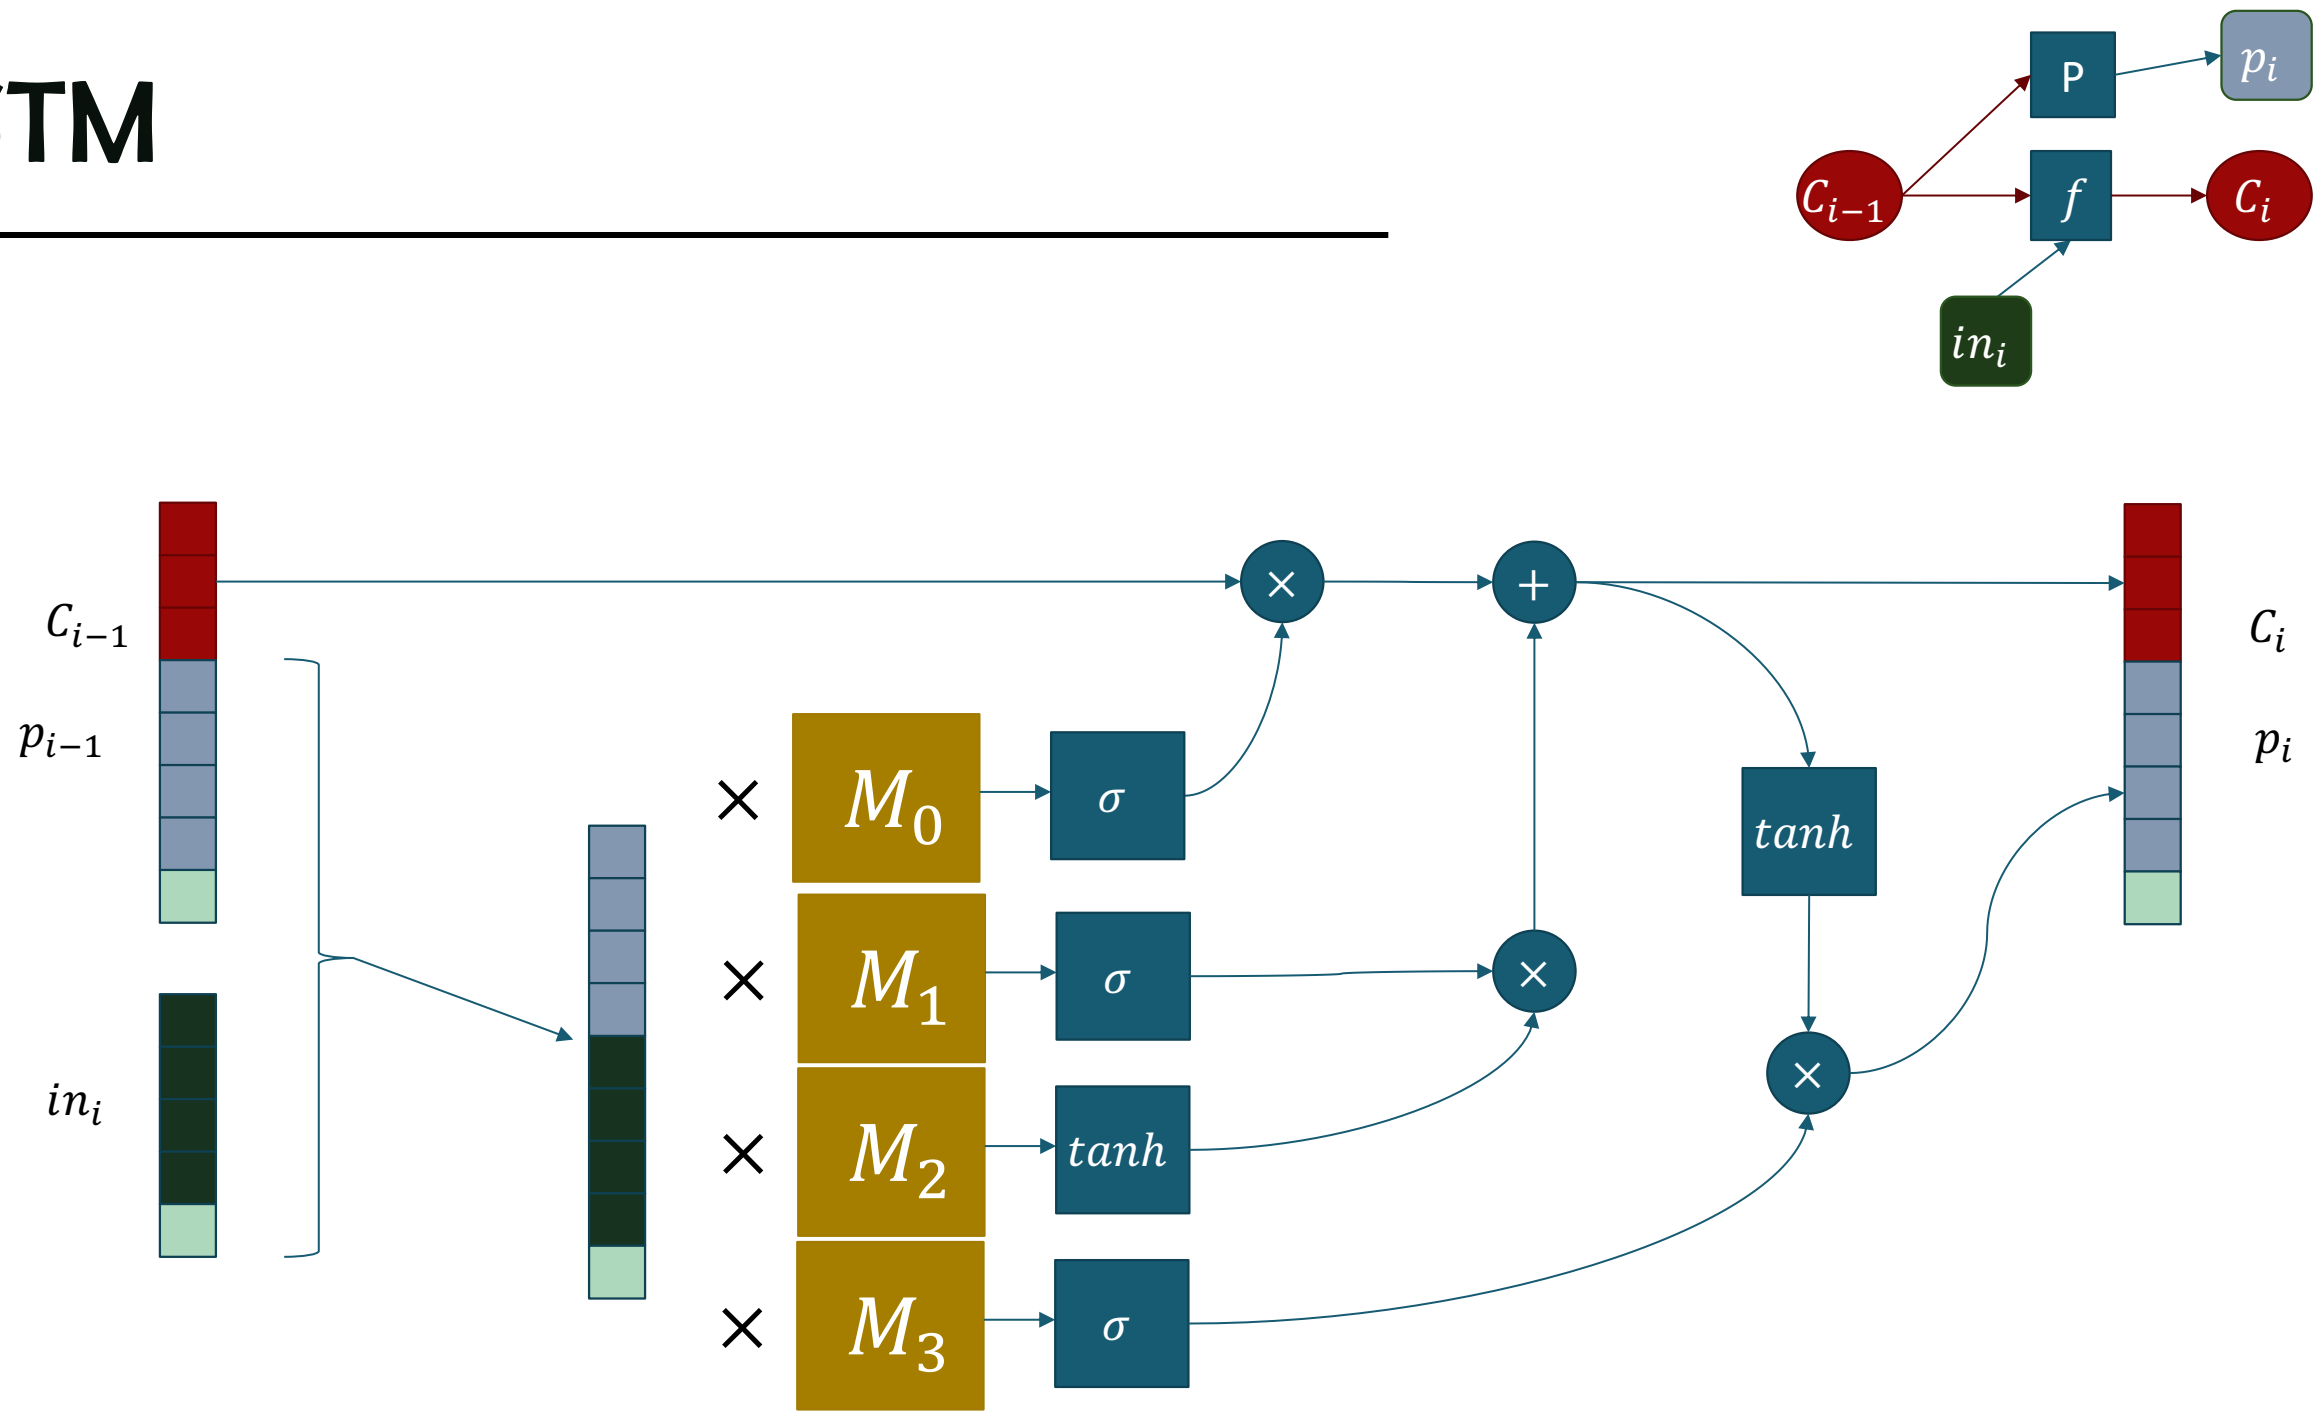


# Simple recurrent neural network

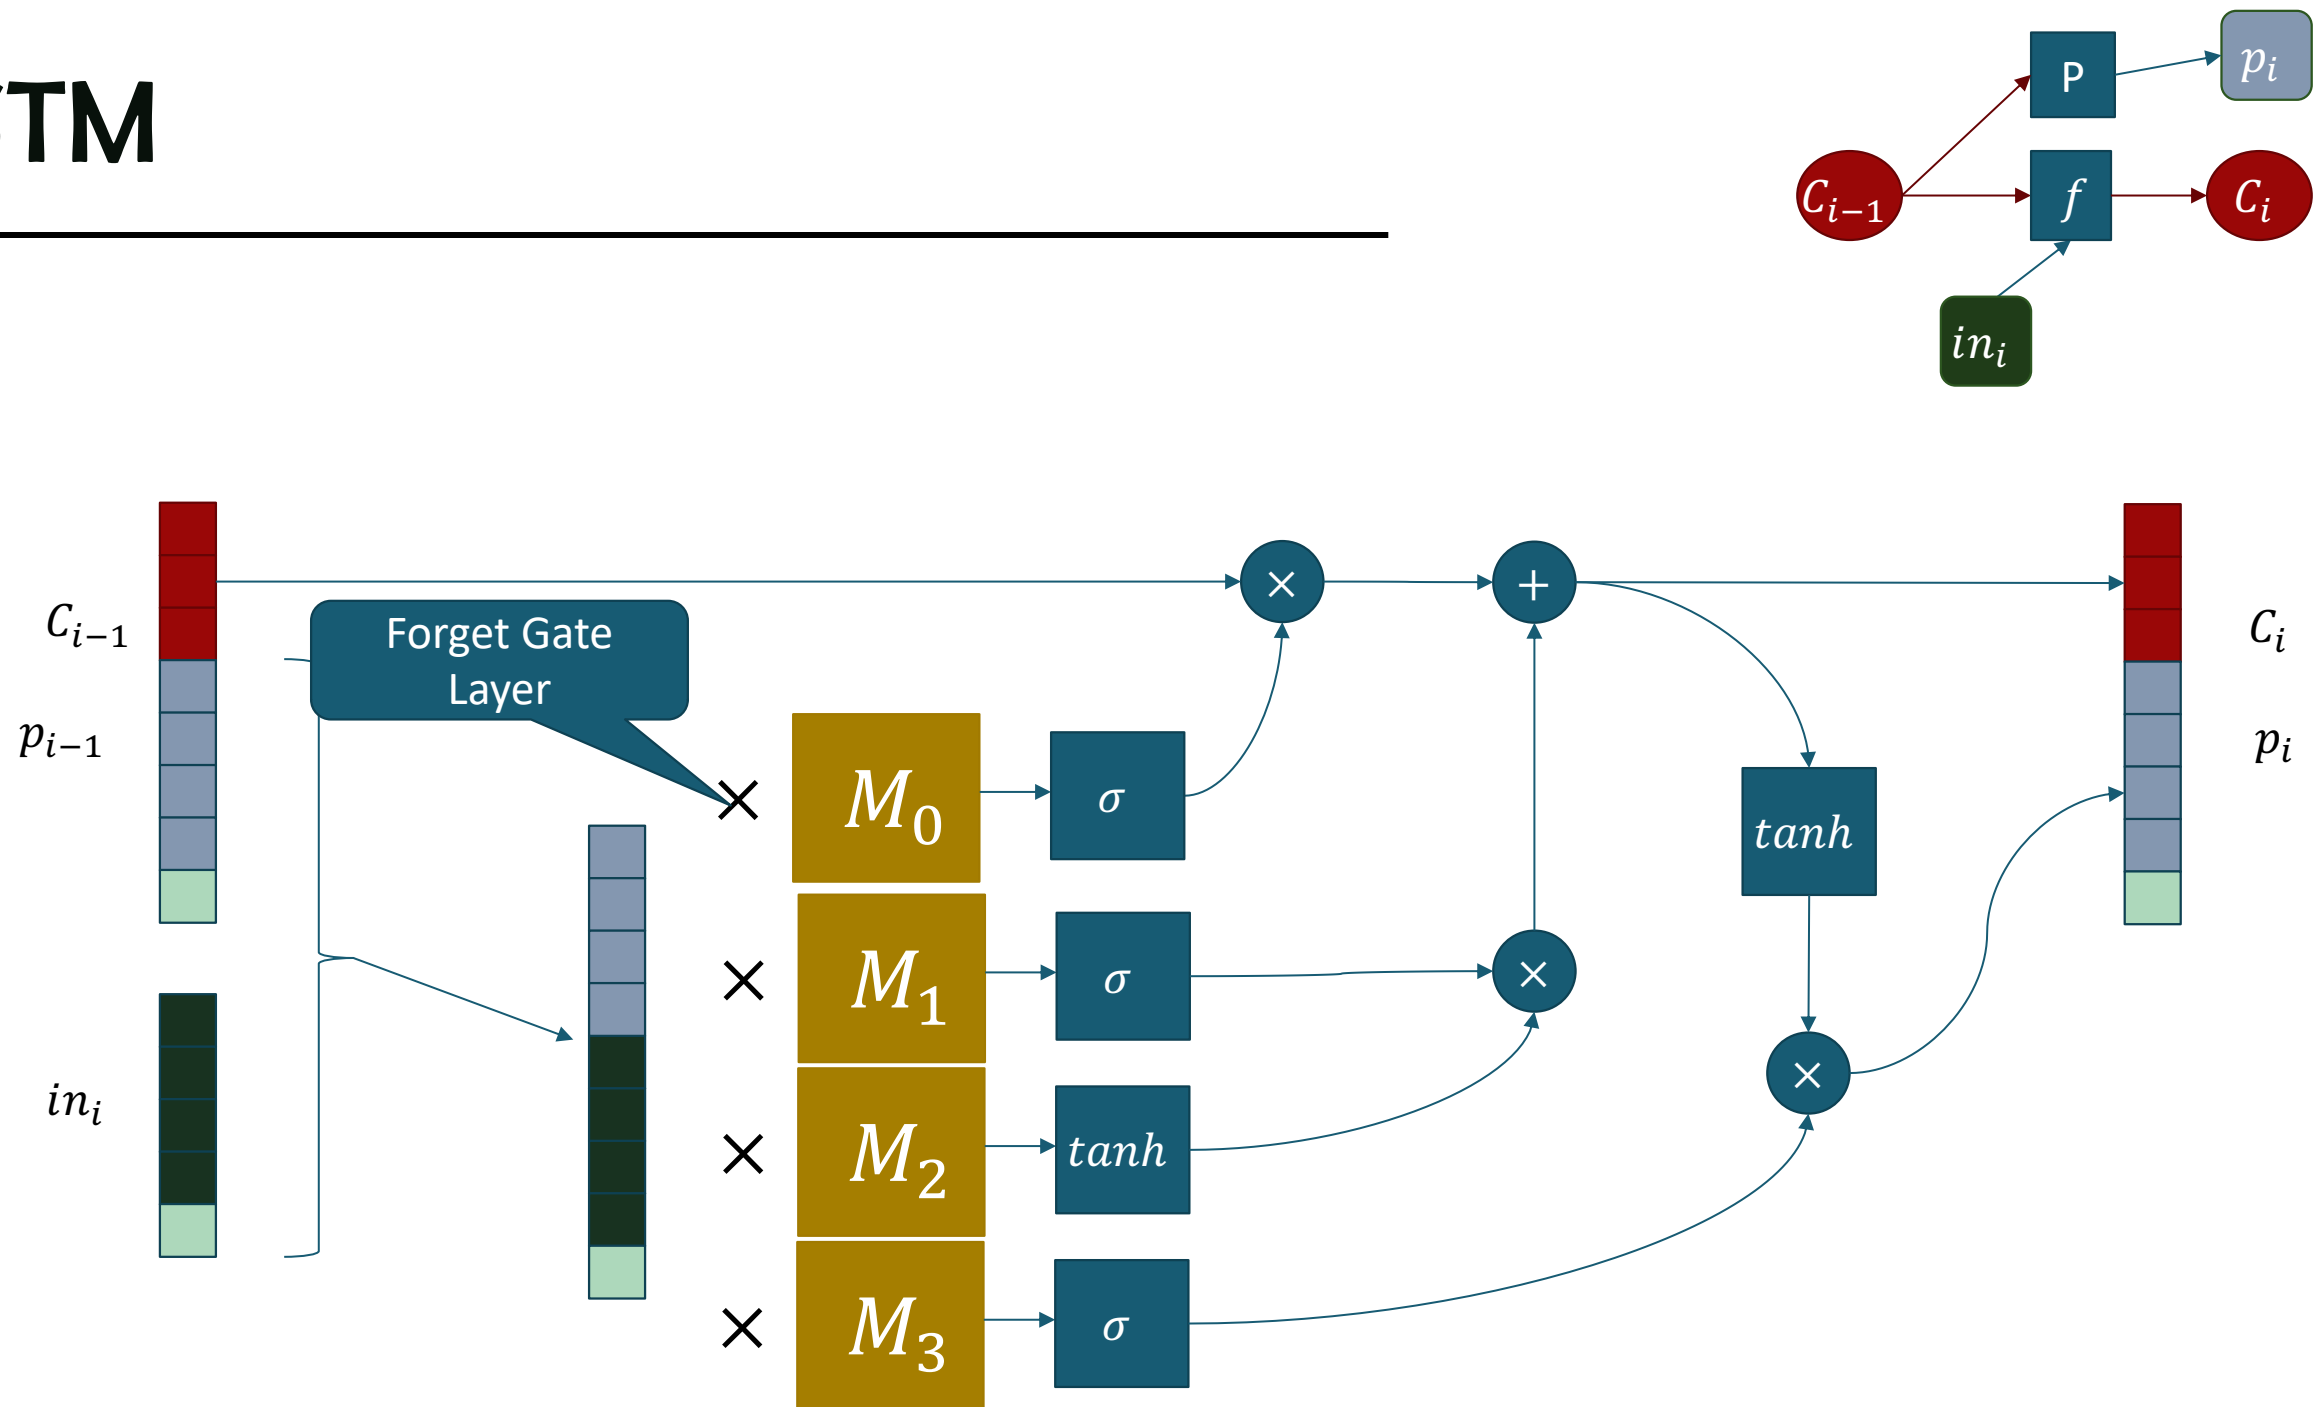
---



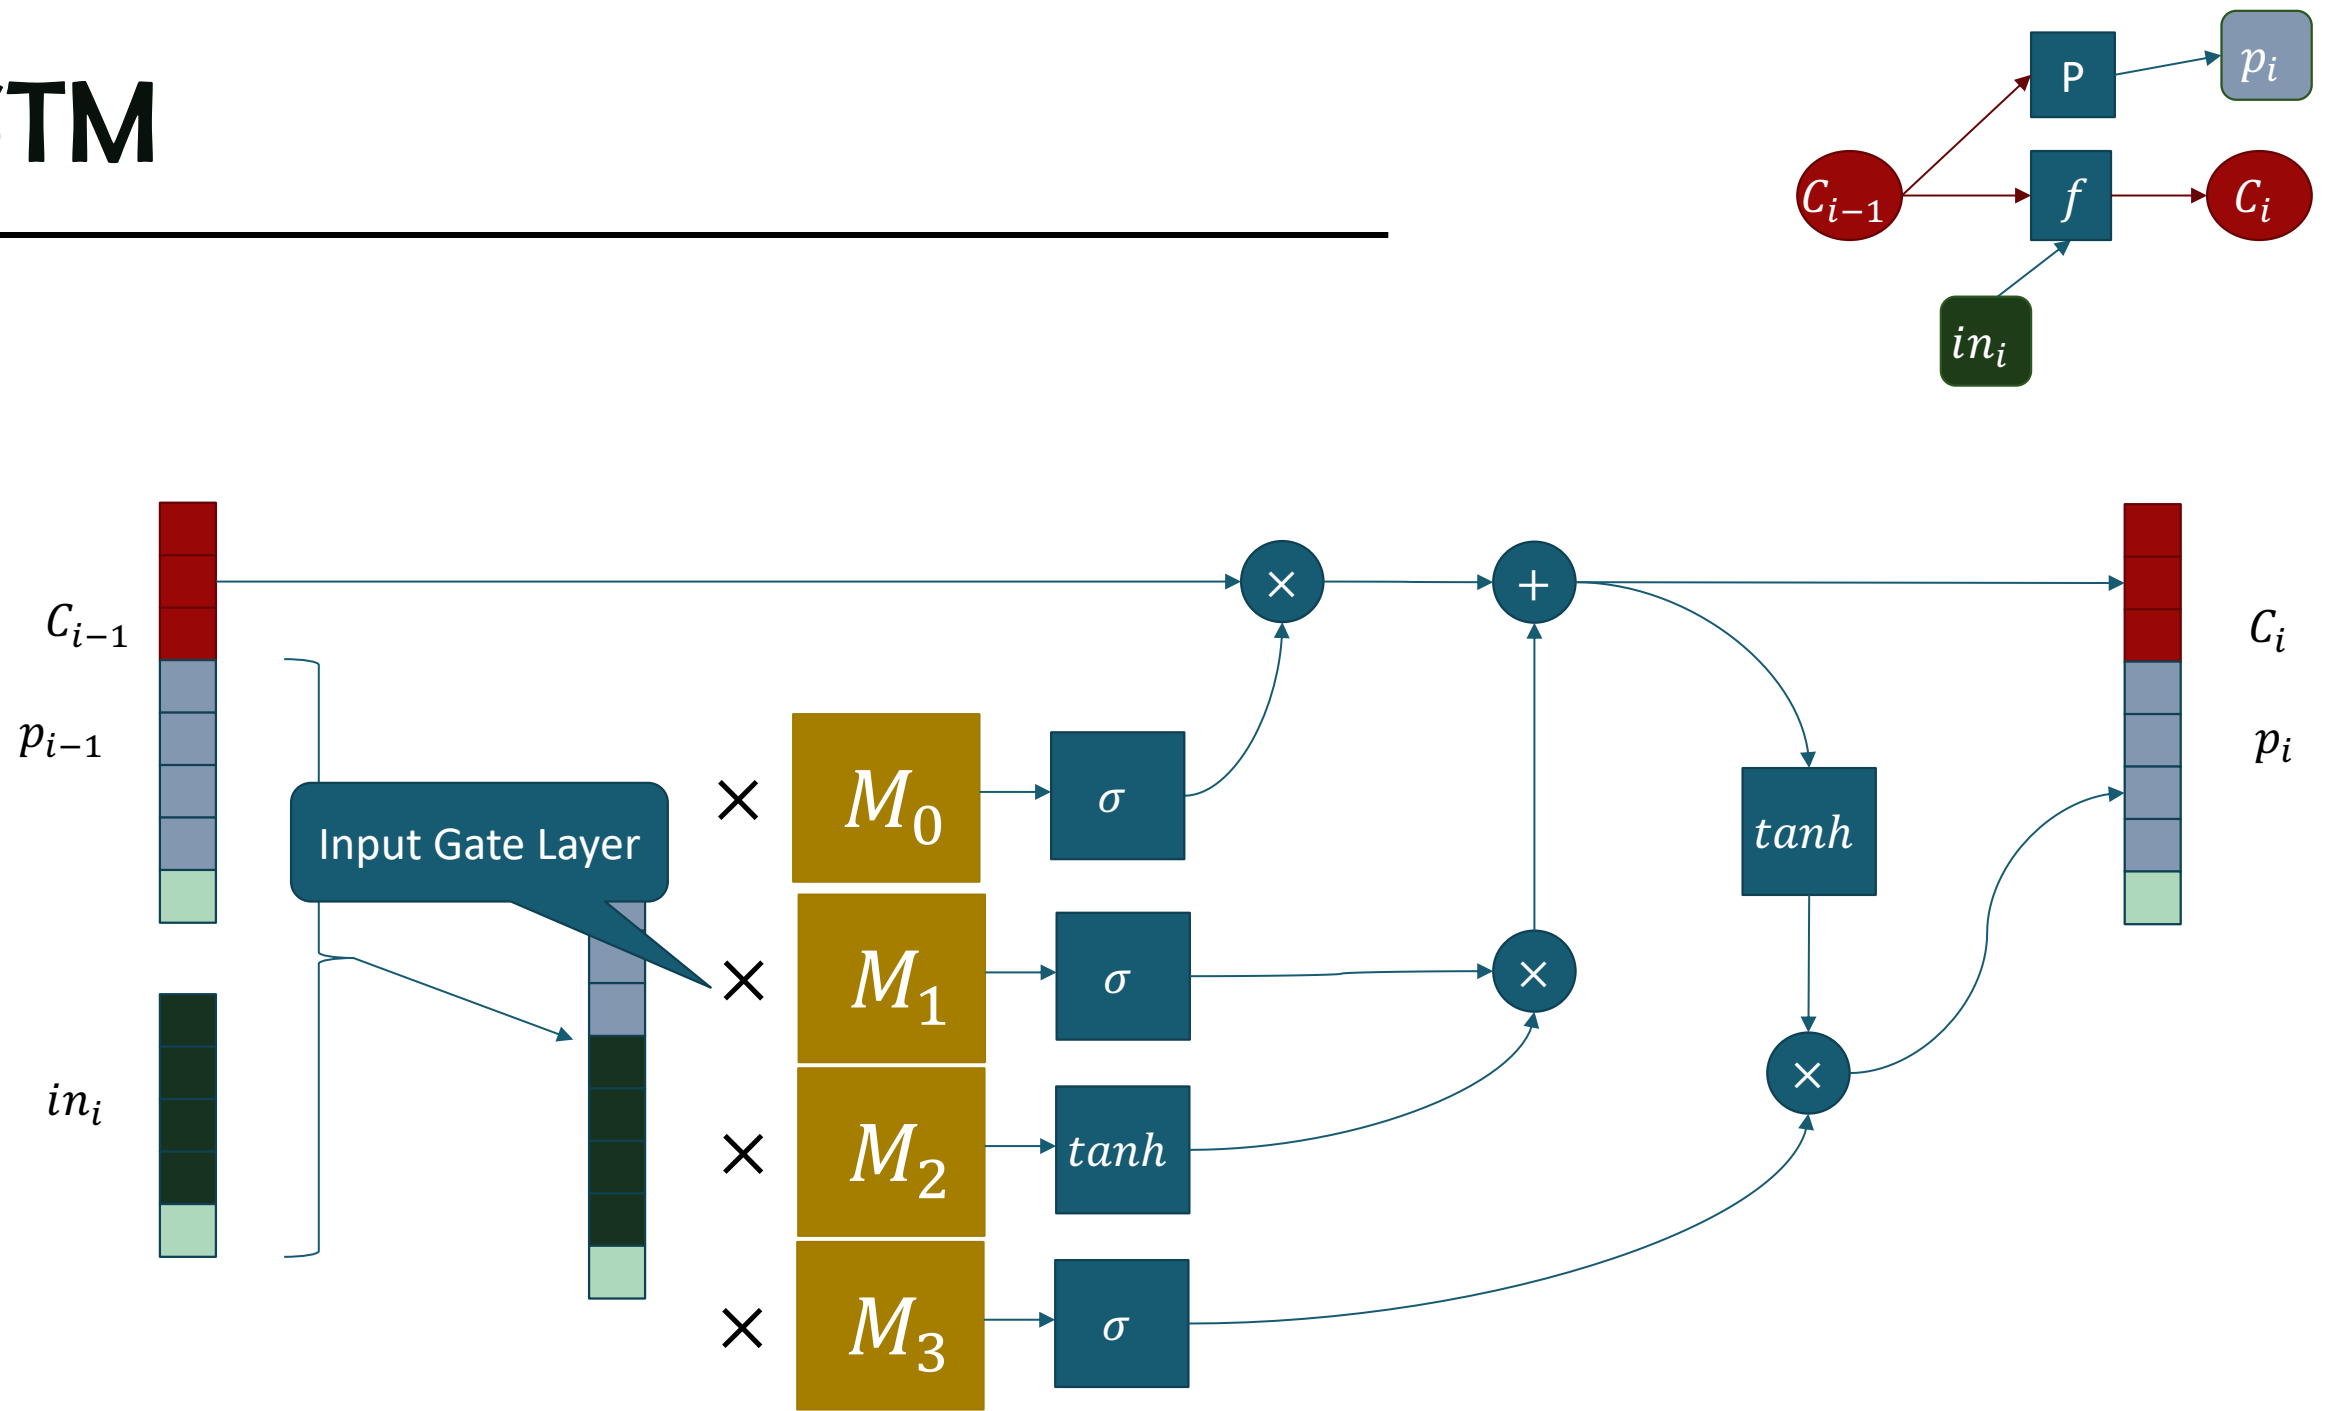
# LSTM



# LSTM



# LSTM







# Encoder-decoder model

Encoder:  
Generates latent  
representation of the evidence

Why do children hate the big brown bear?




Decoder:  
Produces a distribution based  
on the evidence

The big brown bear scares the children with its roar

# Multilayer and bidirectional models

---

Bidirectional networks propagate information back-to-front as well as front-to-back. They can also have more than one layer.




# Attention

---

Key idea: Summarizing into a single vector is a big bottleneck.  
Every output should have direct access to the whole input

# Attention




# RobustFill

---

Single example case

- Attend over the input/output to produce each token

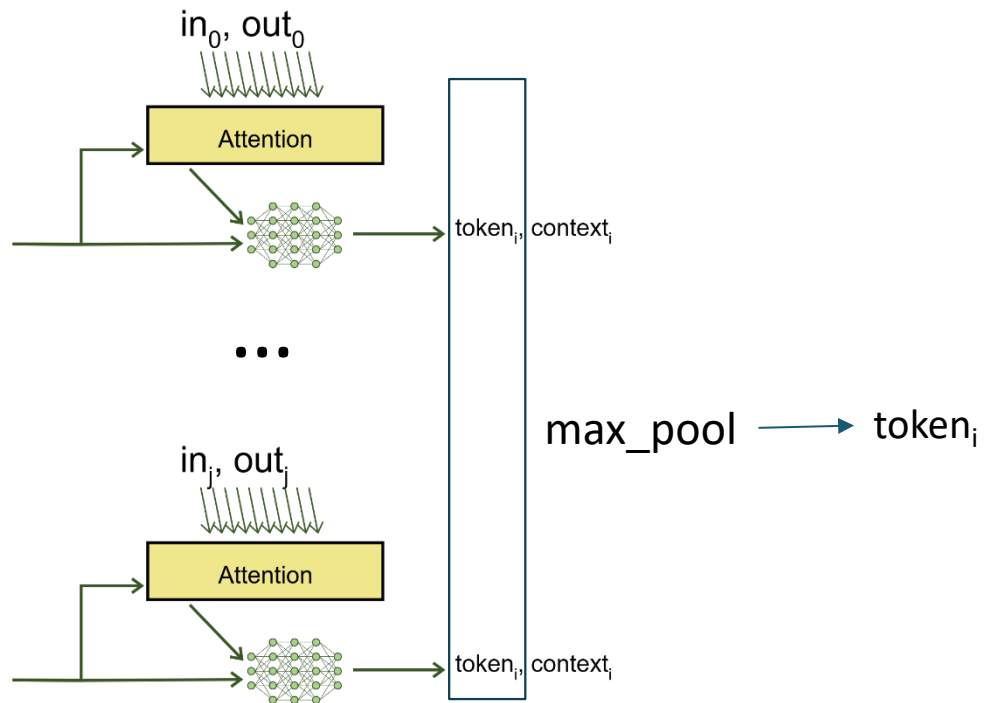


# RobustFill


---

## Multiple example case

- Compute each example independently
- pool over the result to decide which token to output




# Beam Search




At each step pick the K most likely successors

# Beam Search







# Beam Search



# Beam Search



# Beam Search



# Beam search for programs

---

Similar to top-down search

- Same strategies for pruning the search space apply

# Bayesian View of PBE

---

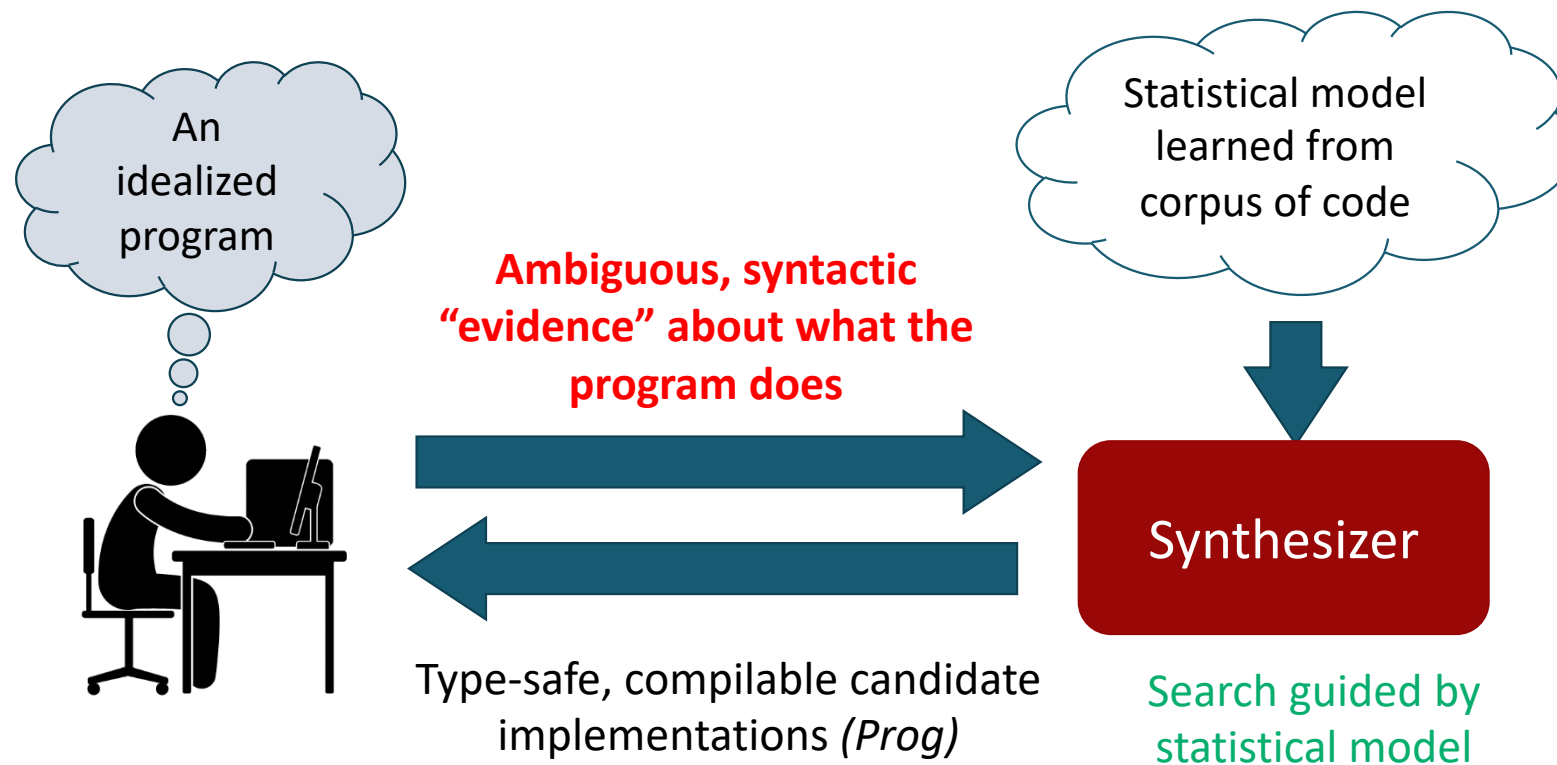
$$P(A | B) = \frac{P(B | A)P(A)}{P(B)}$$

$$P(f | evidence) = \frac{P(evidence | f)P(f)}{P(evidence)}$$

Can we learn  $P(f|evidence)$  directly?

# Synthesis from ambiguous evidence

---



***Neural Sketch Learning for Conditional Program Generation.*** Murali, Qi, Chaudhuri, and Jermaine. ICLR 2018 (oral presentation).

# The Bayou system: Example

---

## Input

```
void read(String file) {  
    /// call:readline  
    /// type:BufferedReader  
}
```




## Output

```
void read(String file) {  
  
    String s;  
    BufferedReader br;  
    FileReader fr;  
    try {  
        fr = new FileReader(file);  
  
        br = new BufferedReader(fr);  
  
        while ((s = br.readLine())!=null){  
        }  
        br.close();  
    }  
    catch (FileNotFoundException _e) {}  
    catch (IOException _e) {}  
}
```

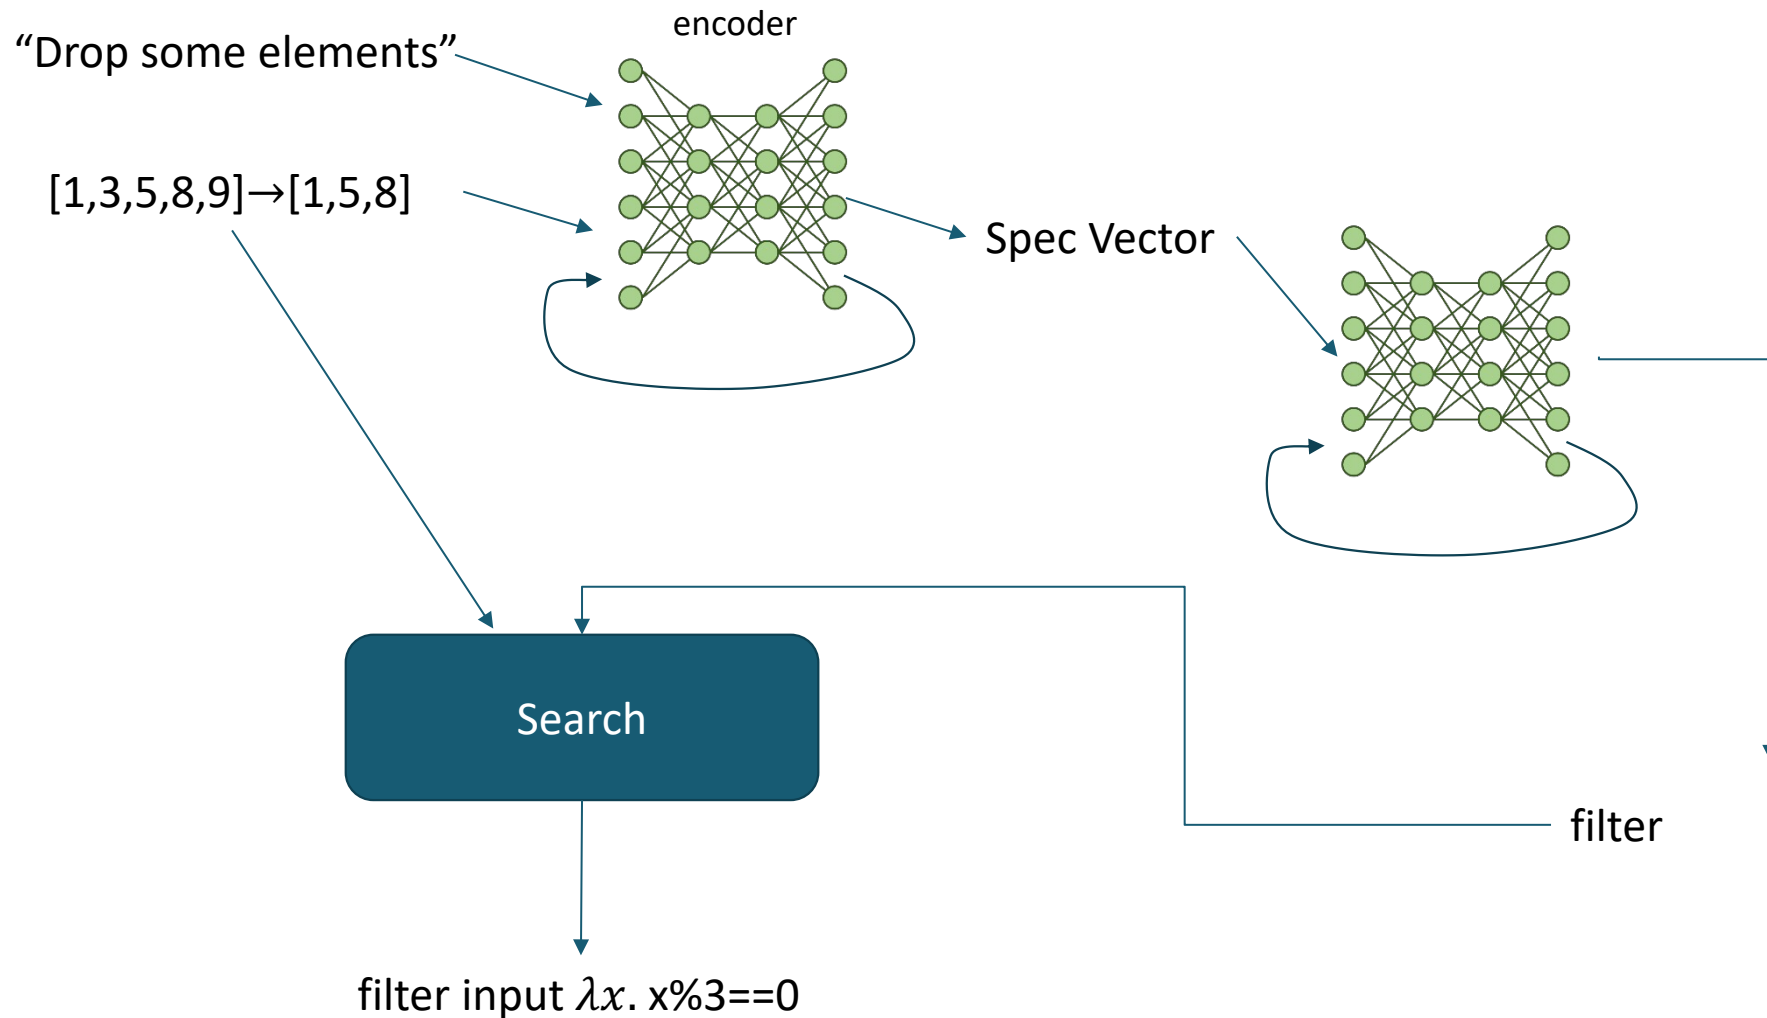
# Learning to Sketch

---






# Learning to Sketch



# Beam Search + Synthesis

---



# Training


---

## Key idea

Penalize heavily for incorrect token  
Penalty for hole depends on cost  
of completing it

“Drop some elements”

$[1,3,5,8,9] \rightarrow [1,5,8]$



filter input  $\lambda x. x \% 3 == 0$

fold input  $\lambda x. x \% 3 == 0$  **X**

filter input **HOLE**  $\lambda x. \text{HOLE}$

# Algolisp Results

---

