

# Lecture 36

# Neurosymbolic Synthesis - I


*Sankha Narayan Guria*

*with slides from*


*Armando Solar-Lezama in-turn based on slides from Swarat Chaudhuri*

# Neurosymbolic program synthesis


## Neural Guided Search




## Symbolic Guided DL




## Distillation



## Relaxation




## Component Discovery



# Neural Guided Search

---




Leverage the ability of NN to learn complex conditional distributions

Network guides the search for programs that satisfy the goals

# Symbolic Guided Deep Learning

---




Symbolic knowledge can be used to guide the training of neural networks

- When you want a network that is consistent with prior knowledge
- When you want to improve data efficiency and better generalization

# Distillation

---



Use the neural network as a starting point for program synthesis

Replace neural components with symbolic ones

- To improve interpretability and analyzability
- To better generalize out of distribution
- To ensure more predictable behavior

# Relaxation

---

$in \rightarrow \lambda \rightarrow out$



$in \rightarrow \lambda \rightarrow out$

Replace symbolic components with neural proxies

- Can help leverage the information in the symbolic component into a larger DL pipeline
- Can help guide the search for symbolic components

**Relaxation:**

**Supporting numerical optimization  
over discrete programs**

# The Parameter Synthesis Problem

---

```
tOff := ??;  tOn := ??;  h := ??  
forever{  
  temp := readTemp();  
  if (isOn() && temp > tOff)  
    switchHeaterOff();  
  elseif ( !isOn() && temp < tOn)  
    switchHeaterOn(h);  
}
```




**Warming:**

$$\frac{d}{dt}temp = -k \cdot temp + h$$

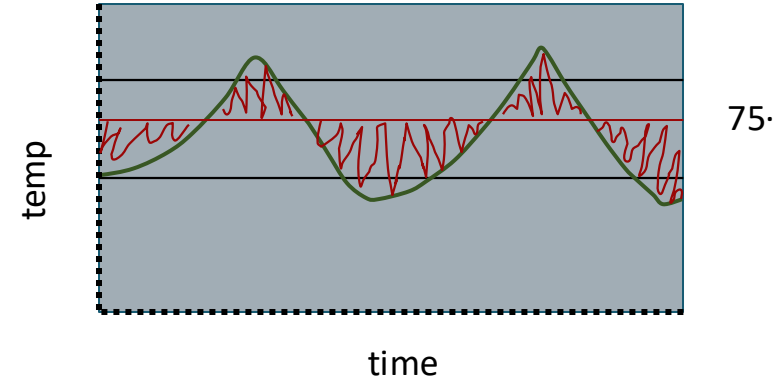

**Cooling:**

$$\frac{d}{dt}temp = -k \cdot temp$$





# The Parameter Synthesis Problem



$$F(tOff, tOn, h) = [t_0, t_1, \dots, t_k]$$

$$Err(tOff, tOn, h) := \sum (t_i - 75)^2$$


Can we find values of  $(tOff, tOn, h)$  to minimize Err?

# Limitations

---

Smooth function converges to original function according to L2 norm.

- That's not enough for optimization.
- The function below converges to the original function, but the local minima is in the worst possible point



# Limitations

---

Programs tend to be difficult to optimize

- Even with smoothing, there are too many local minima


Can be useful in combination with other search techniques

- e.g. Hamiltonian Monte Carlo

**Relaxation + Top-Down = Near**


# Learning to synthesize incrementally 2

---



# Guiding program search

---




So far

- Neural network policy to guide search
- Value function scores individual states

# Guiding program search

---



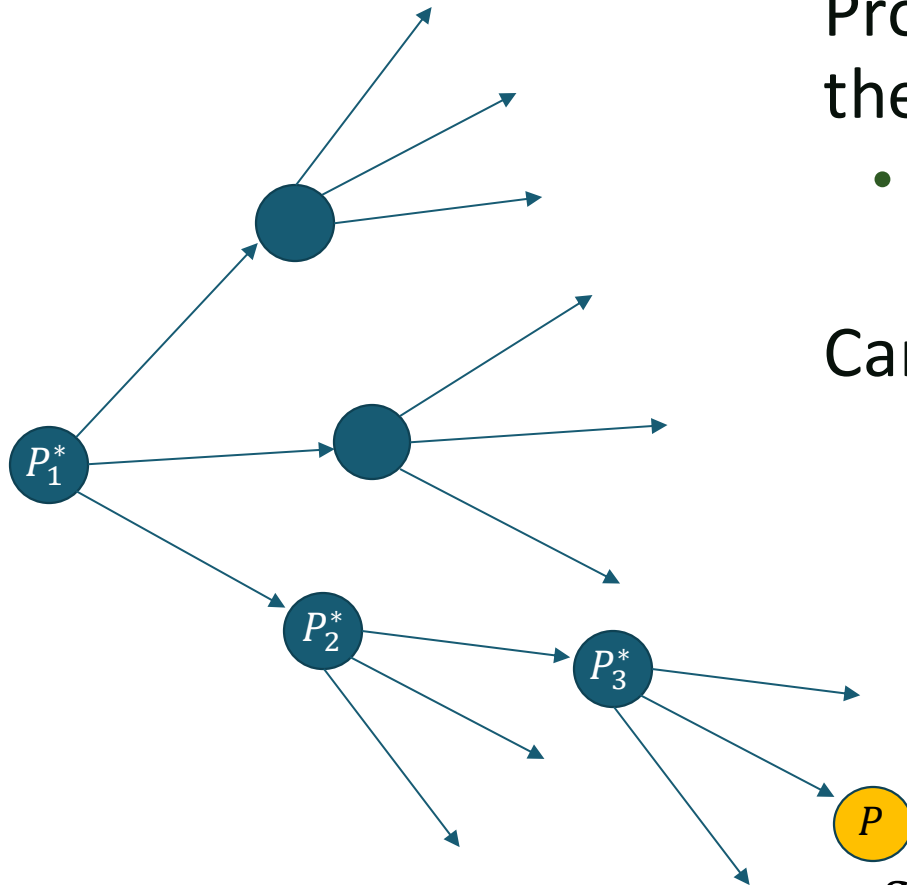
Problem: You only get ground truth on the leaves of the search tree

- Value for an intermediate node is only an estimate

$$s(P) + \min_{\theta} \text{Loss}(\alpha_P, \theta_P)$$

# Guiding program search

---



Problem: You only get ground truth on the leaves of the search tree

- Value for an intermediate node is only an estimate

Can we get a better estimate with DL?

$$s(P) + \min_{\theta} \text{Loss}(\alpha_P, \theta_P)$$



# Estimating the “Cost to Go”

---

$P^*$  = partial program (non-terminal nodes)

$P^*[? \rightarrow e]$  = completions of  $P^*$  (reachable terminal nodes)

$$\text{Heuristic Estimate: } d(P^*) \approx \min_{P \in \mathbb{C}(P^*)} \left[ \underbrace{\Delta s(P, P^*)}_{\text{Additional Structure Cost}} + \underbrace{\min_{\theta} \text{Loss}(\alpha_P, \theta_P)}_{\text{Training Loss}} \right]$$

“Cost to Go”



If  $d(P^*)$  is a lower bound it becomes an “admissible heuristic”

---

# Learning with Neural Heuristics

Shah et al. *Learning Differentiable Programs with Neural Admissible Heuristics*. NeurIPS 2020.

# Guiding Search with Neural Relaxations




**Fill hole with NN**

**Train parameters**

**Use training loss to guide search**


If a large neural network cannot fit this hole, then a neurosymbolic completion also cannot



Motivating Observation/Assumption:

# ~~Functional Representational Power~~

Neurosymbolic Models




“Large” Neural Models

“Neural Relaxation”

Every neurosymbolic model can be (approximately) represented by some “large” neural model.

# Implication (abstract form)

---



Large Neural

Slack due to approximation error or training ability

$$\forall P, \exists f \in F \text{ s.t. } d(f) \leq d(P) + \epsilon$$

Neurosymbolic

Any Cost Function


**We can train an admissible heuristic!**

## “Neural Relaxation”

Every neurosymbolic model can be (approximately) represented by some “large” neural model.

# Informed Search (e.g., A\*)

Use  $d(P^*)$  to prune the search



Can Prune This Branch!

Suppose:

$$s(\lambda x. \text{map } x \, g^*) + d(\lambda x. \text{map } x \, g^*) > s(\lambda x. x) + \text{Loss}(\lambda x. x)$$

Structural Cost

Training Loss

"Cost to Go" Heuristic

# A\* Search

---

Priority queue of current leaf nodes:

- Sorted by  $s(P^*) + d(P^*)$

Pop off top program  $P^*$

- If  $P^*$  is complete, terminate
- Else, expand  $P^*$ , add child nodes to priority queue


Lower bounds “Cost to Go”

**Guarantee:** if  $d(P^*)$  is **admissible**, A\* will return optimal  $P$

- Tighter  $d(P^*)$  prunes more aggressively
- Uninformed  $d(P^*)$  (e.g., always 0) => uninformed search



# NEAR: Results

Order of magnitude speedup!





# NEAR: Neural Admissible Relaxations



**Fill hole with NN**

**Train parameters**

**Use training loss as admissible heuristic**

If a large neural network cannot fit this hole, then a neurosymbolic completion also cannot

